

A decentralized approach to publishing confidential data

Master Thesis Presentation 14.07.2009

Konstantin Welke



Professorship for Communication Systems
Department of Computer Science
Albert Ludwigs University of Freiburg

Outline

- 1 Motivation
- 2 Cryptography
- 3 The Protocol
- 4 Conclusion

Outline

- 1 Motivation
- 2 Cryptography
- 3 The Protocol
- 4 Conclusion

Motivation



- This is Bob
- Bob hosted a party at his house yesterday
- He made some pictures to share with his friends
- How can he do that?

Motivation



- This is Alice
- Alice maintains an electronic calendar
- Some appointments should be visible
 - Only to her co-workers
 - Only to her friends
 - Only to her family
 - Only to herself
- How can she do that?

Motivation

- How can Bob share his pictures with his friends?
- How can Alice share appointments with defined groups of peers?
- **But no one else!**
 - 1 Social Networks / Web Applications?
 - 2 Just set up a server?
 - 3 Upload an encrypted tar archive?

The Confidential Publishing Problem

- More general: How can people share content with peers over third-party infrastructure, with
 - Confidentiality
 - Integrity
 - Authenticity
 - Decentralized structure
 - Usability
- How to achieve that?

Outline

- 1 Motivation
- 2 **Cryptography**
- 3 The Protocol
- 4 Conclusion

How to solve the confidential publishing problem

- Confidentiality
 - Exchange symmetric key using public-key crypto
 - Encrypt using symmetric crypto
 - Problem: How to encrypt for n people?
- Integrity & Authenticity
 - Hash using one-way hash function
 - Sign using public-key crypto

How to solve the confidential publishing problem

- Decentralized structure
 - Just like “traditional internet services”
 - user@server
- Usability
 - Easy to use & to try out
 - Web application
 - User friendly crypto?
 - Problem: How to verify public-key fingerprints?

How to encrypt for n people?


- Crypto aims at one-to-one encryption
- How to encrypt for n people?
- Naive approach
 - Generate symmetric key k
 - Share k with set of people U
 - If user u joins:
 - Give k to u
 - Constant cost
 - If user u is removed:
 - Generate new key k'
 - Share k' with $U \setminus u$
 - Linear cost

How to encrypt for n people?

- Can we do better than linear removal cost?
- Idea: Build a tree (Wallner et al. [1999], Wong et al. [2000])


Key Trees

- User knows exactly the keys from their leaf key to the group key




Adding a user

- Create new leaf for the user, send leaf to user
- If tree is full: create new group key




Adding a user

- Create new leaf for the user, send leaf to user
- If tree is full: create new group key




Removing a user

- All keys that the removed user knows need to be exchanged with new keys



Removing a user

- All keys that the removed user knows need to be exchanged with new keys



Efficiency

- Naive approach
 - Constant cost to add
 - Linear cost to remove
- Key Tree
 - Amortized constant cost to add
 - Logarithmic cost to remove

Public-key Fingerprints

- Users need to verify public-key fingerprints to detect man-in-the-middle-attacks
 - Real-life example: SSH

- Are these the same fingerprint?

`c1:b1:30:29:d7:b8:de:6c:97:77:10:d7:46:41:63:87`

`c1:b1:30:29:d7:b8:de:6c:77:97:10:d7:46:41:63:87`

- User friendly?

Human-Readable Fingerprints

- Introduced by Bååth and Kühn [2002], Kaminsky [2006] independently
- Generate human-readable sentences from the fingerprint
- Are these the same fingerprint?

Happy:

Lisa:chased:the cookie:

Savah:jumped over:the flower:

Chip:smiled at:the world:

Nacho:joked about:the pizza

Happy:

Lisa:chased:the cookie:

Savah:jumped over:the flower:

Chip:smiled at:the pizza:

Nacho:joked about:the world

Outline

- 1 Motivation
- 2 Cryptography
- 3 The Protocol**
- 4 Conclusion

The Protocol

- User should be able to publish content to a set of peers with
 - Confidentiality
 - Integrity
 - Authenticity
 - Decentralized structure
 - Usability
- We covered the cryptography that is needed
- Let's describe the protocol!

Login

- Every user has a passphrase
- Generate symmetric key from passphrase
- Private content is encrypted with that key
 - E.g. the user's private key

Short-term passphrases

- User has *one main* passphrase
 - Encrypts his long-term public key
 - Valid for e.g. 5 years
- Problem: Should the user enter that passphrase at an Internet cafe?
- Solution: *Short-term* passphrases
 - Valid for a short period of time
 - Encrypt a *short-term public key*
 - Can only access a subset of data
 - Only that data can be compromised
 - Usable?

Public keys

- Folder “/public-keys”
 - Contains all public keys
 - Publicly accessible
- Public keys have validity period
- *Key messages*
 - Signed by preceeding key:
 - Introduce new keys
 - Declare old keys compromised
 - *Not* signed:
 - Declare old keys “lost”

Content Files

- Content is organized into folders and files
- *Content files* consists of
 - Actual (encrypted) content
 - Revision number
 - Signature
 - Key to access the content
 - Encrypted for every user/group that may access the content
- One content file can consist of multiple filesystem files
 - E.g. one containing content, another containing the signature
- Folders are just a “list of files”

File names

- Problem: Files names carry potentially confidential information
 - “Bob totally drunk”
 - “Alice’s doctor appointment”
- Solution: Generate *public filename(s)* from *private name*
 - Every folder defines a *salt*
 - Public name: Hash of salt + user id + namespace + private name
 - Namespace: e.g. content, revision, ...
 - If content file is split over multiple filesystem files
 - Avoids name clashes

File names

- Example:

- Salt: "ssssssssssssssssssssssssssssssssss"
- Public name: "Alice's doctor appointment"
- Namespace: "content"
- Private name:

SHA-256("ssssssssssssssssssssssssssssssssss"
+ "alice@crypto.uni-freiburg.de:"
+ "content:/")
+ "Alice's doctor appointment") =

1afac6c112b4b5a1a9f43ff8445fdebc43e00c41078edfbad3b6d533e5efc15f

- Hard to guess private name from public name


Publishing Content

- Users organize their contacts into groups
 - Every group = one key tree
- When publishing a content file, user can set which users/groups have read access


Social Network?

- *Privacy-respecting* social network becomes possible
- User maintain groups of contacts
 - Only visible to the group
- User publishes content
 - Select target audience
 - Not visible to service provider
 - Or any third party
- Also possible:
 - Messaging
 - “Wall”
 - Discussion groups
 - Tagging


Example




Example




Example



Example



Example



Outline

- 1 Motivation
- 2 Cryptography
- 3 The Protocol
- 4 Conclusion

Conclusion

- Protocol solves confidential publishing problem
- Users can publish data over third-party infrastructure
 - Without revealing the data to the infrastructure
 - One-to-many encryption is crucial
- Web application prototype
 - Incomplete (only single-user)

That's all folks!

Thank you for your attention!

Do you have any questions?

... & Demo